



The ARM Instruction Set

Main features of the ARM Instruction Set

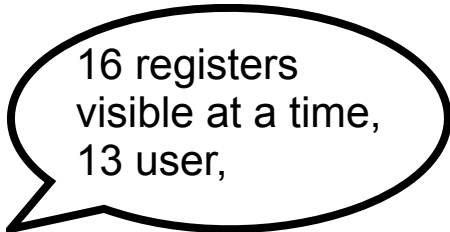
- * **All instructions are 32 bits long.**
 - * **Most instructions execute in a single cycle.**
 - * **Every instruction can be conditionally executed.**
 - * **A load/store architecture**
 - Data processing instructions act only on registers
 - Three operand format
 - Combined ALU and shifter for high speed bit manipulation
 - Specific memory access instructions with powerful auto-indexing addressing modes.
 - 32 bit and 8 bit data types
and also 16 bit data types on ARM Architecture v4.
 - Flexible multiple register load and store instructions
 - * **Instruction set extension via coprocessors**
- * 16 registers (32 bits)
 - * PC is a regular register
=> can read/write with regular instructions

The Registers

- * **ARM has 37 registers in total, all of which are 32-bits long.**
 - 1 dedicated program counter
 - 1 dedicated current program status register
 - 5 dedicated saved program status registers
 - 30 general purpose registers
- * **However these are arranged into several banks, with the accessible bank being governed by the processor mode. Each mode can access**
 - a particular set of r0-r12 registers
 - a particular r13 (the stack pointer) and r14 (link register)
 - r15 (the program counter)
 - cpsr (the current program status register)

and privileged modes can also access

 - a particular spsr (saved program status register)

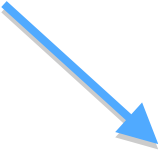


16 registers
visible at a time,
13 user,

Register Organisation

General registers and Program Counter

User mode
RF view



User32 / System	FIQ32	Supervisor32	Abort32	IRQ32	Undefined32
r0	r0	r0	r0	r0	r0
r1	r1	r1	r1	r1	r1
r2	r2	r2	r2	r2	r2
r3	r3	r3	r3	r3	r3
r4	r4	r4	r4	r4	r4
r5	r5	r5	r5	r5	r5
r6	r6	r6	r6	r6	r6
r7	r7	r7	r7	r7	r7
r8	r8_fiq	r8	r8	r8	r8
r9	r9_fiq	r9	r9	r9	r9
r10	r10_fiq	r10	r10	r10	r10
r11	r11_fiq	r11	r11	r11	r11
r12	r12_fiq	r12	r12	r12	r12
r13 (sp)	r13_fiq	r13_svc	r13_abt	r13_irq	r13_undef
r14 (lr)	r14_fiq	r14_svc	r14_abt	r14_irq	r14_undef
r15 (pc)	r15 (pc)	r15 (pc)	r15 (pc)	r15 (pc)	r15 (pc)

Program Status Registers

cpsr	cpsr	cpsr	cpsr	cpsr	cpsr
	spsr_fiq	spsr_svc	spsr_abt	spsr_irq	spsr_undef

Accessing Registers using ARM Instructions

- * **No breakdown of currently accessible registers.**
 - All instructions can access r0-r14 directly.
 - Most instructions also allow use of the PC.
- * **Specific instructions to allow access to CPSR and SPSR.**
- * **Note : When in a privileged mode, it is also possible to load / store the (banked out) user mode registers to or from memory.**
 - See later for details.

The Program Status Registers (CPSR and SPSRs)



Copies of the ALU status flags (latched if the instruction has the "S" bit set).

- * **Condition Code Flags**

N = **N**egative result from ALU flag.

Z = **Z**ero result from ALU flag.

C = ALU operation **C**arried out

V = ALU operation o**V**erflowed

- * **Mode Bits**

M[4:0] define the processor mode.

- * **Interrupt Disable bits.**

I = 1, disables the IRQ.

F = 1, disables the FIQ.

- * **T Bit (Architecture v4T only)**

T = 0, Processor in ARM state

T = 1, Processor in Thumb state

Condition Flags

	Logical Instruction	Arithmetic Instruction
<u>Flag</u>		
Negative (N='1')	No meaning	Bit 31 of the result has been set Indicates a negative number in signed operations
Zero (Z='1')	Result is all zeroes	Result of operation was zero
Carry (C='1')	After Shift operation '1' was left in carry flag	Result was greater than 32 bits
oVerflow (V='1')	No meaning	Result was greater than 31 bits Indicates a possible corruption of the sign bit in signed numbers

The Program Counter (R15)

- * **When the processor is executing in ARM state:**
 - All instructions are 32 bits in length
 - All instructions must be word aligned
 - Therefore the PC value is stored in bits [31:2] with bits [1:0] equal to zero (as instruction cannot be halfword or byte aligned).
- * **R14 is used as the subroutine link register (LR) and stores the return address when Branch with Link operations are performed, calculated from the PC.**
- * **Thus to return from a linked branch**
 - `MOV r15,r14`

or

 - `MOV pc,lr`

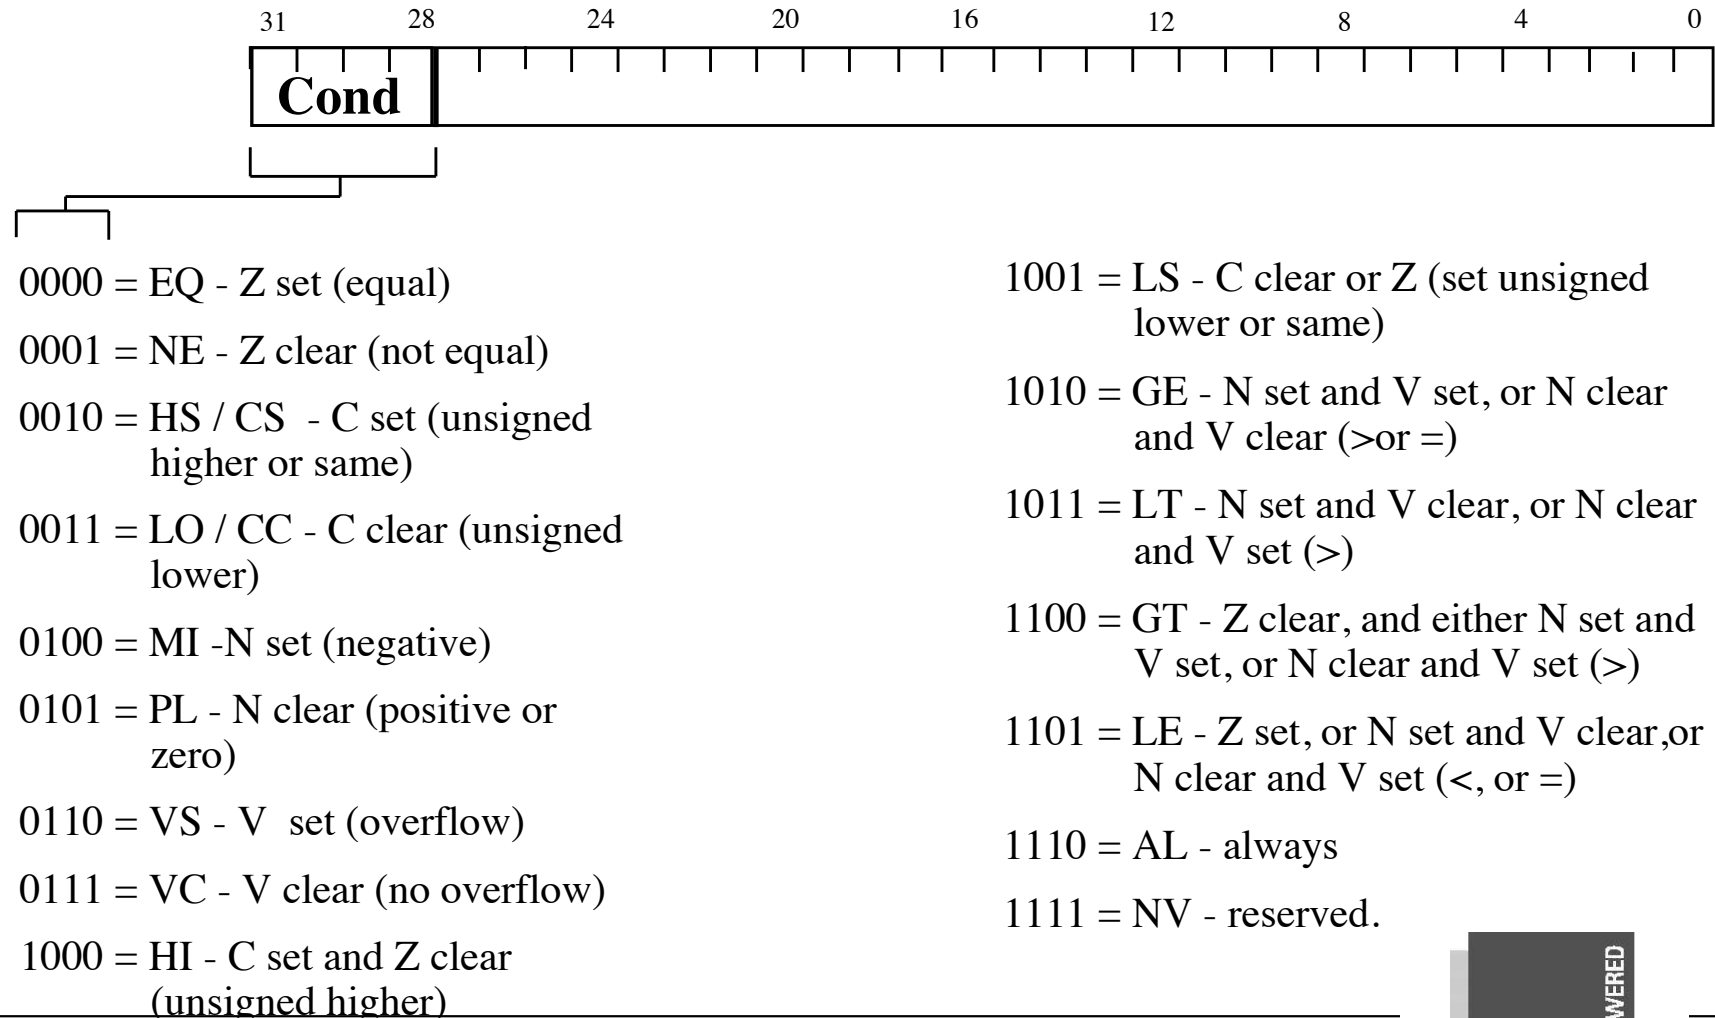
ARM Instruction Set Format

31	2827				1615				87				0				<u>Instruction type</u>												
Cond		0	0	I	Opcode			S	Rn		Rd		Operand2				Data processing / PSR Transfer												
Cond		0	0	0	0	0	0	0	A	S	Rd		Rn		Rs	1	0	0	1	Rm	Multiply								
Cond		0	0	0	0	0	1	U	A	S	RdHi		RdLo		Rs	1	0	0	1	Rm	Long Multiply (v3M / v4 only)								
Cond		0	0	0	1	0	B	0	0		Rn		Rd		0	0	0	0	1	0	0	1	Rm	Swap					
Cond		0	1	I	P	U	B	W	L		Rn		Rd		Offset						Load/Store Byte/Word								
Cond		1	0	0	P	U	S	W	L		Rn		Register List										Load/Store Multiple						
Cond		0	0	0	P	U	1	W	L		Rn		Rd		Offset1	1	S	H	1	Offset2			Halfword transfer : Immediate offset (v4 only)						
Cond		0	0	0	P	U	0	W	L		Rn		Rd		0	0	0	0	1	S	H	1	Rm	Halfword transfer: Register offset (v4 only)					
Cond		1	0	1	L	Offset															Branch								
Cond		0	0	0	1	0			0	1	0	1			1	1	1	1			1	1	1	0	0	0	1	Rn	Branch Exchange (v4T only)
Cond		1	1	0	P	U	N	W	L		Rn		CRd		CPNum		Offset						Coprocessor data transfer						
Cond		1	1	1	0	Op1					CRn		CRd		CPNum		Op2		0	CRm			Coprocessor data operation						
Cond		1	1	1	0	Op1				L	CRn		Rd		CPNum		Op2		1	CRm			Coprocessor register transfer						
Cond		1	1	1	1	SWI Number															Software interrupt								

Conditional Execution

- * **Most instruction sets only allow branches to be executed conditionally.**
- * **However by reusing the condition evaluation hardware, ARM effectively increases number of instructions.**
 - All instructions contain a condition field which determines whether the CPU will execute them.
 - Non-executed instructions soak up 1 cycle.
 - Still have to complete cycle so as to allow fetching and decoding of following instructions.
- * **This removes the need for many branches, which stall the pipeline (3 cycles to refill).**
 - Allows very dense in-line code, without branches.
 - The Time penalty of not executing several conditional instructions is frequently less than overhead of the branch or subroutine call that would otherwise be needed.

The Condition Field




Using and updating the Condition Field

- * **To execute an instruction conditionally, simply postfix it with the appropriate condition:**
 - For example an add instruction takes the form:
 - `ADD r0,r1,r2 ; r0 = r1 + r2 (ADDAL)`
 - To execute this only if the zero flag is set:
 - `ADDEQ r0,r1,r2 ; If zero flag set then...`
`; ... r0 = r1 + r2`
- * **By default, data processing operations do not affect the condition flags (apart from the comparisons where this is the only effect). To cause the condition flags to be updated, the S bit of the instruction needs to be set by postfixing the instruction (and any condition code) with an “S”.**
 - For example to add two numbers and set the condition flags:
 - `ADDS r0,r1,r2 ; r0 = r1 + r2`
`; ... and set flags`

Branch instructions (1)

- * **Branch :** `B{<cond>} label`
- * **Branch with Link :** `BL{<cond>} sub_routine_label`



- * **The offset for branch instructions is calculated by the assembler:**
 - By taking the difference between the branch instruction and the target address minus 8 (to allow for the pipeline).
 - This gives a 26 bit offset which is right shifted 2 bits (as the bottom two bits are always zero as instructions are word – aligned) and stored into the instruction encoding.
 - This gives a range of ± 32 Mbytes.

Branch instructions (2)

- * **When executing the instruction, the processor:**
 - shifts the offset left two bits, sign extends it to 32 bits, and adds it to PC.
- * **Execution then continues from the new PC, once the pipeline has been refilled.**
- * **The "Branch with link" instruction implements a subroutine call by writing PC-4 into the LR of the current bank.**
 - i.e. the address of the next instruction following the branch with link (allowing for the pipeline).
- * **To return from subroutine, simply need to restore the PC from the LR:**
 - `MOV pc, lr`
 - Again, pipeline has to refill before execution continues.
- * **The "Branch" instruction does not affect LR.**
- * **Note: Architecture 4T offers a further ARM branch instruction, BX**
 - See Thumb Instruction Set Module for details.

Data processing Instructions

- * **Largest family of ARM instructions, all sharing the same instruction format.**
- * **Contains:**
 - Arithmetic operations
 - Comparisons (no results - just set condition codes)
 - Logical operations
 - Data movement between registers
- * **Remember, this is a load / store architecture**
 - These instruction only work on registers, *NOT* memory.
- * **They each perform a specific operation on one or two operands.**
 - First operand always a register - Rn
 - Second operand sent to the ALU via barrel shifter.
- * **We will examine the barrel shifter shortly.**

Arithmetic Operations

* Operations are:

- ADD operand1 + operand2
- ADC operand1 + operand2 + carry
- SUB operand1 - operand2
- SBC operand1 - operand2 + carry - 1
- RSB operand2 - operand1
- RSC operand2 - operand1 + carry - 1

* Syntax:

- <Operation>{<cond>}{S} Rd, Rn, Operand2

* Examples

- ADD r0, r1, r2
- SUBGT r3, r3, #1
- RSBLES r4, r5, #5

Data Movement

* **Operations are:**

- MOV operand2
- MVN NOT operand2

Note that these make no use of operand1.

* **Syntax:**

- <Operation>{<cond>}{S} Rd, Operand2

* **Examples:**

- MOV r0, r1
- MOVS r2, #10
- MVNEQ r1, #0

Loading full 32 bit constants

- * Although the MOV/MVN mechanism will load a large range of constants into a register, sometimes this mechanism will not generate the required constant.
- * Therefore, the assembler also provides a method which will load *ANY* 32 bit constant:
 - `LDR rd,=numeric constant`
- * If the constant can be constructed using either a MOV or MVN then this will be the instruction actually generated.
- * Otherwise, the assembler will produce an LDR instruction with a PC-relative address to read the constant from a literal pool.
 - `LDR r0,=0x42 ; generates MOV r0,#0x42`
 - `LDR r0,=0x55555555 ; generate LDR r0,[pc, offset to lit pool]`
- * As this mechanism will always generate the best instruction for a given case, it is the recommended way of loading constants.

Load / Store Instructions

- * **The ARM is a Load / Store Architecture:**
 - Does not support memory to memory data processing operations.
 - Must move data values into registers before using them.
- * **This might sound inefficient, but in practice isn't:**
 - Load data values from memory into registers.
 - Process data in registers using a number of data processing instructions which are not slowed down by memory access.
 - Store results from registers out to memory.
- * **The ARM has three sets of instructions which interact with main memory. These are:**
 - Single register data transfer (LDR / STR).
 - Block data transfer (LDM/STM).
 - Single Data Swap (SWP).


Single register data transfer

- * **The basic load and store instructions are:**
 - Load and Store Word or Byte
 - LDR / STR / LDRB / STRB
- * **ARM Architecture Version 4 also adds support for halfwords and signed data.**
 - Load and Store Halfword
 - LDRH / STRH
 - Load Signed Byte or Halfword - load value and sign extend it to 32 bits.
 - LDRSB / LDRSH
- * **All of these instructions can be conditionally executed by inserting the appropriate condition code after STR / LDR.**
 - e.g. LDREQB
- * **Syntax:**
 - `<LDR|STR>{<cond>}{<size>} Rd, <address>`

Load and Store Word or Byte: Base Register

* The memory location to be accessed is held in a base register

- STR r0, [r1] ; Store contents of r0 to location pointed to
; by contents of r1.
- LDR r2, [r1] ; Load r2 with contents of memory location
; pointed to by contents of r1.



Load and Store Word or Byte: Offsets from the Base Register

- * As well as accessing the actual location contained in the base register, these instructions can access a location offset from the base register pointer.
- * This offset can be
 - An unsigned 12bit immediate value (ie 0 - 4095 bytes).
 - A register, optionally shifted by an immediate value
- * This can be either added or subtracted from the base register:
 - Prefix the offset value or register with '+' (default) or '-'.
- * This offset can be applied:
 - before the transfer is made: *Pre-indexed addressing*
 - optionally *auto-incrementing* the base register, by postfixing the instruction with an '!'.
 - after the transfer is made: *Post-indexed addressing*
 - causing the base register to be *auto-incremented*.

Coprocessors

- * **The ARM architecture supports 16 coprocessors**
- * **Each coprocessor instruction set occupies part of the ARM instruction set.**
- * **There are three types of coprocessor instruction**
 - Coprocessor data processing
 - Coprocessor (to/from ARM) register transfers
 - Coprocessor memory transfers (load and store to/from memory)
- * **Assembler macros can be used to transform custom coprocessor mnemonics into the generic mnemonics understood by the processor.**
- * **A coprocessor may be implemented**
 - in hardware
 - in software (via the undefined instruction exception)
 - in both (common cases in hardware, the rest in software)